Your support agent is fixable, and it wasn’t the model

From: Arindam Mandal · Partner delivery team
To: Priya Anand · VP, Customer Operations · Meridian

Hi Priya,

Short answer: **it’s fixable, and the model wasn’t the problem.** The instructions around it were.

What was going wrong

The agent’s instructions told it that every ticket gets exactly one specialist, and that once that specialist reports back, the ticket is “handled” and should be closed as resolved. They also told it that escalating to a human is a last resort that counts against the team.

So when a customer writes in with two problems, the agent picks one, fixes it, and closes the ticket anyway. On T-4471 the customer had a company-wide login outage *and* a $1,200 overcharge. The agent fixed the login problem, only acknowledged the refund, and marked the ticket resolved. It did what it was told.

What we changed

One file, the coordinator’s instructions. No model change and no code change.

1. List every problem in the ticket and bring in a specialist for each one, most urgent first.
2. Check each specialist’s result before replying. A problem only counts as handled if it was actually fixed.
3. Only claim what a specialist really did. Close a ticket as resolved only when everything is handled; otherwise escalate or ask the customer for what’s missing.
4. Send SSO and Okta login problems to the account specialist, even when the error message looks technical. This was a second, smaller routing mistake we found along the way.

The proof

Same ticket (T-4471), same model, run five times each:

	Resolved	Cost per ticket
Before	0 of 5	$0.13
After	5 of 5	$0.21

We also tested three tickets we hadn’t tuned against. The first version of the fix resolved 8 of 9 of those runs. We haven’t re-run them since the last routing change, and five runs is a small sample, so treat 5 of 5 as strong evidence, not a guarantee. Cost per ticket rises

because a ticket with two problems now uses two specialists.

“Why not just use a better model?”

We tried. We ran the same ticket on a larger, more expensive model with the original instructions. It still resolved 0 of 5, at about twice the cost per ticket ($0.25 vs $0.13). The instructions hadn’t changed, so the one-specialist-then-close rule was still there, and the bigger model didn’t get around it. Rewriting the instructions took the cheaper model from 0 of 5 to 5 of 5. A bigger model would have cost more and fixed nothing.

What’s left

The fix itself is small and already tested. What remains is rollout work:

- Re-run the held-out tickets after the last change.
- Build a larger set of real multi-issue tickets from your history to test against.
- Keep a check running that confirms “resolved” means everything was handled, as the prompts evolve.

The constraint to hold is cost per ticket, which should stay around $0.21.

Happy to walk your leadership through the numbers whenever it’s useful.

Best,
Arindam